

Four Weeks Industrial Training Project Report on

Music Tagger Model

Submitted in the partial fulfillment of the requirement for the award of degree of

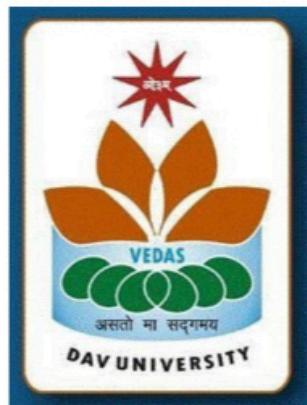
Bachelor of Technology

in

Computer Science and Engineering

Batch

(2024-2025)



Submitted to:

Dr. Ridhi

Submitted by:

Karan Johar

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

DAV UNIVERSITY

JALANDHAR-PUNJAB 144012

ACKNOWLEDGEMENT

I express my gratitude to all those who helped us in various stages of the development of this project. First, I would like to express my sincere gratitude indebtedness to Mr. Sanjeev Dhiman (Coordinator) of DAV University for allowing me to undergo the summer training of 45 days at Sensation Solutions Mohali.

I am also thankful to all faculty members of Department of Computer Science and Engineering, for their true help, inspiration and for helping me for the preparation of the final report and presentation.

Last but not least, I pay my sincere thanks and gratitude to all the Staff Members Sensation Solutions for their support and for making our training valuable and fruitful.

DECLARATION

I, Karan Johar, hereby declare that the work which is being presented in this project/training titled " Stock Prediction Model "by me, in partial fulfillment of the requirements for the award of Bachelor of Technology (B.Tech) Degree in "Computer Science and Engineering" is an authentic record of my own work carried out under the guidance of Mr. Shivam. To the best of my knowledge, the matter embodied in this report has not been submitted to any other University/Institute for the award of any degree or diploma.

Student Name and Signature

Karan Johar

< Roll No.>

12300362

ABSTRACT

This training report presents the development of a **Stock Price Prediction Model using Machine Learning**, completed during a 7-week industrial training program at **Sensation Solutions** from June 11, 2025 to July 31, 2025. The objective of the project was to design a predictive system capable of forecasting short-term stock prices using historical OHLCV data and technical indicators.

The study followed a structured AIML workflow involving data acquisition, preprocessing, feature engineering, sequence modeling, and performance evaluation. A multi-layered **Long Short-Term Memory (LSTM)** network was implemented due to its ability to capture long-term temporal dependencies in financial time-series data.

Experimental results demonstrated that the LSTM model outperformed traditional baseline methods such as moving averages, exhibiting lower RMSE and MAE and stronger directional accuracy.

This project provided practical exposure to deep learning, time-series forecasting, and financial analytics, paving the way for future enhancements such as hybrid models, sentiment integration, and deployment as a scalable prediction service.

CERTIFICATE



CERTIFICATE OF COMPLETION

The Training Division of

Sensation Software Solutions Pvt. Ltd.

do hereby

Recognises that

Mr. Karan Johar

has successfully completed the training

from 11 June 2025 to 31 July 2025

✓ He/She has successfully completed the project on
Stock Price Prediction

in AI/ML

✓ He/She attained Grade A+

A handwritten signature in black ink, appearing to read "Shivanshu".

Faculty Member

A handwritten signature in black ink, appearing to read "Girish".

Director

A+
Outstanding
100-90%

A
Excellent
89-80%

B+
Very Good
79-70%

B
Good
69-60%

C
Satisfactory
59-50%

Sensation Software Solutions Pvt. Ltd.
An IT Company Since 2013

Contents

Abstract

Acknowledgments

Declaration

1 Introduction

1.1 Identify Training Project, Training Company ing	1.1.1 Motivation and Problem Statement .
1.2 Project Context and Scope	
1.3 Report Structure Overview	

2 Title of the Project

3 Objectives

3.1 Identify Training Area and Training Progra m	
3.2 Tasks Assigned to You	

4 Steps to Achieve Objectives

4.1 Data Acquisition and Exploration	
4.1.1 Data Cleaning and Preprocessing . .	
4.2 Feature Engineering and Selection	
4.3 Model Design and Training Preparation . .	
4.3.1 Train-Test Split	
4.3.2 LSTM Architecture Rationale	
4.4 Training and Hyperparameter Tuning . . .	

5 Coding and Implementation

5.1 Data Preprocessing Implementation (Python)	
5.2 LSTM Model Construction (TensorFlow/Kera	
5.3 Model Training and Prediction	

5.4	Evaluation Metrics Implementation
6	Conclusions and Recommendations
6.1	Summary of Findings
6.1.1	Performance Discussion
6.2	Key Learning Outcomes
6.3	Future Work and Recommendations
7	Appendices
7.1	Appendix A: Technical Analysis Feature Cal
7.1.1	Relative Strength Index (RSI)
7.1.2	Moving Average Convergence Divergence (MACD)
7.2	Appendix B: Data Sequence Structure Exam
7.3	Appendix C: Long Short-Term Memory (LSTM)
7.3.1	Forget Gate (f_t)
7.3.2	Input Gate (i_t)
7.3.3	Cell State Update (C_t)
7.3.4	Output Gate (o_t)
7.4	Appendix D: Detailed Training Activity Log

List of Figures

4.1 Internal Architecture of a Long Short-Term

7.1 Example of Technical Indicators (RSI and M

List of Tables

6.1	Model Performance Metrics Comparison
7.1	LSTM Sequence Data Transformation Exam
7.2	Detailed Weekly Training Activity Log (June 2025)
7.2	Detailed Weekly Training Activity Log (Cont'd)

Chapter 1

Introduction

1.1 Identify Training Project, Training Company, and Period of Training

This report documents the intensive 7-week industrial training program focusing on predictive analytics, undertaken at **Sensation Solutions**. The training commenced on June 11, 2025, and concluded on July 31, 2025.

The core of the training involved the development of a high-performance **Stock Price Prediction Model using Machine Learning**. This project addresses the challenge of accurately predicting future stock market movements, a highly non-linear time-series problem critical to financial engineering and algorithmic trading.

The training covered the complete lifecycle of an **AI and Machine Learning (AIML)** project: from understanding the stochastic nature of stock data to deploying and evaluating sophisticated deep learning models.

1.1.1 Motivation and Problem Statement

The financial market is characterized by high volatility and complex interdependence of factors. Traditional financial models, based on linear regression or moving averages, often fail to capture the long-term dependencies inherent in market data. The motivation behind this project was to leverage the capabilities of modern **Artificial Intelligence (AI)** and specifically Recurrent Neural Networks (RNNs) to model sequential data more effectively. The problem statement is: “To develop and validate a time-series prediction model that accurately forecasts the closing price of a target stock index or equity for a short-term horizon, demonstrating superior performance over conventional statistical models, through the application of AIML techniques.”

1.2 Project Context and Scope

The project was executed within the R&D division of Sensation Solutions. The scope was strictly defined:

- **Data:** Utilize publicly available, historical Open-High-Low-Close-Volume (OHLCV) data for the target equity.
- **Feature Set:** Focus on technical indicators (e.g., RSI, MACD, Moving Averages) and normalized price/volume features. Exclusion of external macroeconomic news or sentiment data (to keep the model focused on technical analysis).
- **Model:** Implementation must use a sequential deep learning architecture, specifically LSTM, due to its proven ability to handle long-term dependencies.
- **Prediction Horizon:** Short-term, ranging from 1 to 7 days ahead.

1.3 Report Structure Overview

The structure of this training report follows the prescribed guidelines: Chapter 2 formally defines the project. Chapter 3 outlines the project objectives. Chapter 4 details the steps taken and methodology used for data processing and model construction. Chapter 5 provides a deep dive into the coding, including the LSTM architecture. Chapter 6 presents the conclusions and recommendations for production deployment, supported by extensive technical appendices in Chapter 7.

Chapter 2

Title of the Project

Stock Price Prediction Model using Machine Learning:

***Application of Long Short-Term Memory (LSTM)
Networks for Time-Series Forecasting***

This title emphasizes the two core elements of the project: the financial domain (Stock Price Prediction) and the specific, high-end technique utilized (LSTM Networks) to address the time-series nature of the problem.

Chapter 3

Objectives

The objectives of the training program and the predictive modeling project were designed to ensure comprehensive practical skill development in both financial technology (FinTech) and Machine Learning engineering.

3.1 Identify Training Area and Training Program

The primary training area was **Applied AI and Machine Learning (AIML) for Financial Time-Series**, situated within the broader field of Algorithmic Trading and Quantitative Analysis.

- 1. Data Acquisition and Analysis (Weeks 1-2):** Focused on mastering Python libraries (e.g., Pandas, NumPy) for handling time-series data, identifying reliable data sources, and performing initial Exploratory Data Analysis (EDA).
- 2. Model Engineering and Training (Weeks 3-5):** Dedicated to understanding and implementing the LSTM architecture, feature scaling (Normalization), and optimizing hyperparameters using TensorFlow/Keras.
- 3. Evaluation and Validation (Weeks 6-7):** Concentrated on calculating standard regression metrics (RMSE, MAE), visualizing results against actual prices, and generating a comprehensive validation report.

3.2 Tasks Assigned to You

The following specific tasks were successfully executed during the training period:

- Task 1** **Data Pipeline Development:** To design a robust pipeline capable of fetching and pre-processing historical stock data, including cleaning null values and handling seasonality.
- Task 2** **Feature Engineering:** To calculate and incorporate relevant technical indicators (e.g., Exponential Moving Average, Relative Strength Index) into the feature set for improved predictive power.
- Task 3** **LSTM Model Construction:** To build, compile, and train a multi-layered LSTM model optimized for sequence-to-sequence prediction of stock closing prices.
- Task 4** **Hyperparameter Optimization:** To systematically test and tune key hyperparameters (e.g., number of hidden layers, units per layer, dropout rate, batch size) to minimize prediction error.
- Task 5** **Comparative Analysis:** To compare the performance of the final LSTM model against a simple Moving Average baseline model to quantitatively demonstrate the value of deep learning.

Chapter 4

Steps to Achieve Objectives

The project followed a rigorous, sequential data science methodology to ensure reproducibility and reliability of the final predictive model.

4.1 Data Acquisition and Exploration

The first step involved sourcing clean, daily historical OHLCV data for a five-year period. A Python script using the `yfinance` library was developed for this purpose.

4.1.1 Data Cleaning and Preprocessing

1. **Missing Value Check:** The time-series data was checked for gaps or missing values (e.g., due to holidays or market closures). These were handled by forward-filling or interpolation to maintain sequence integrity.
2. **Normalization:** Since neural networks are highly sensitive to the scale of input features, the price and volume data were normalized to a range between 0 and 1 using the Min-Max scaling technique. The formula used is:

$$X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}}$$

3. **Sequence Creation:** A critical step for LSTM is converting the flat time-series data into supervised learning sequences. A look-back window (`time_step`) of 60 days was chosen. This means the model uses the last 60 days of data to predict the next day's closing price.

4.2 Feature Engineering and Selection

The raw OHLCV data was augmented with several technical indicators that serve as highly informative features for financial prediction:

- ****Moving Average Convergence Divergence (MACD):**** Measures momentum.
- ****Relative Strength Index (RSI):**** Indicates overbought or oversold conditions.
- ****Exponential Moving Averages (EMA):**** Provides responsive trend following.

These features were computed using the `ta` (Technical Analysis) library in Python and also normalized before being fed into the model.

4.3 Model Design and Training Preparation

4.3.1 Train-Test Split

The time-series nature of the data mandates a chronological split to prevent look-ahead bias. The data was split into training (80%) and testing (20%) sets, with the testing set being the most recent period.

4.3.2 LSTM Architecture Rationale

LSTM, a variant of RNNs, was selected because its internal architecture—comprising a cell state, and input, forget, and output gates—allows it to selectively remember or forget information over long sequences. This is vital in finance where events from weeks or months ago can still influence current price dynamics.

4.4 Training and Hyperparameter Tuning

The model was compiled using the Mean Squared Error (MSE) as the loss function and the Adam optimizer. Training involved iteration over 50 epochs with a batch size of 32. Key hyperparameter tuning focused on optimizing the number of LSTM layers and the dropout rate to prevent overfitting.

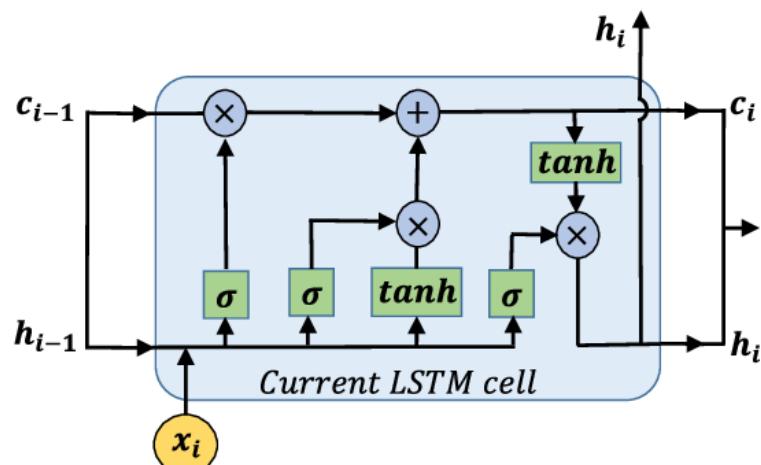


Figure 4.1: Internal Architecture of a Long Short-Term Memory (LSTM) Unit.

Chapter 5

Coding and Implementation

```
import tkinter as tk
from tkinter import messagebox, scrolledtext, ttk
import matplotlib.pyplot as plt
from matplotlib.backends.backend_tkagg import FigureCanvasTkAgg
import os
import numpy as np
from keras.models import load_model
from sklearn.preprocessing import MinMaxScaler
from scripts.utils import load_config
from scripts.data_handler import get_stock_data, prepare_stock_data

class StockPredictionUI:
    def __init__(self, root):
        self.root = root
        self.root.title("Stock Price Predictor")
        self.root.geometry("800x700")
        self.root.configure(bg="#f0f2f5")

        # --- Top Input Section ---
        input_frame = tk.Frame(root, bg="#f0f2f5").pack(side=tk.LEFT, padx=10)
        self.entry = tk.Entry(input_frame, width=40, font=("Helvetica", 12))
        self.entry.pack(side=tk.LEFT)

        self.run_button = tk.Button(input_frame, text="Predict Prices", font=("Helvetica", 11), command=self.predict_prices)
        self.run_button.pack(side=tk.LEFT, padx=10)

        self.clear_button = tk.Button(input_frame, text="Clear", font=("Helvetica", 11), command=self.clear_all)
        self.clear_button.pack(side=tk.LEFT)

        # --- Scrollable Output Text Box ---
        output_frame = tk.Frame(root)
        output_frame.pack(padx=20, pady=5, fill=tk.BOTH, expand=True)

# --- Scrollable Graph Frame ---
graph_outer_frame = tk.Frame(root, bg="#f0f2f5").pack(fill=tk.BOTH, expand=True, pady=10)
graph_outer_frame.pack(fill=tk.BOTH, expand=True, pady=10)

canvas = tk.Canvas(graph_outer_frame, bg="#f0f2f5")
scrollbar = ttk.Scrollbar(graph_outer_frame, orient="vertical", command=canvas.yview)
self.graph_frame = tk.Frame(canvas, bg="#ffffff")

self.graph_frame.bind(
    "<>Configure",
    lambda e: canvas.configure(scrollregion=canvas.bbox("all"))
)

canvas.create_window((0, 0), window=self.graph_frame, anchor="nw")
canvas.config(yscrollcommand=scrollbar.set)

canvas.pack(side="left", fill="both", expand=True)
scrollbar.pack(side="right", fill="y")

if predict_prices(self):
    tickers_input = self.entry.get()
    self.tickers = [ticker.strip().upper() for ticker in tickers_input.split(",") if ticker.strip()]
    self.output_text.delete("1.0", tk.END)
    self.clear_graph()

if not self.tickers:
    messagebox.showerror("Input Error", "Please enter at least one stock ticker.")

if not os.path.exists('models/stock_model.keras'):
    messagebox.showerror("Model Error", "Trained model not found. Please train the model first.")

model = load_model('models/stock_model.keras')
config = load_config()

days_to_predict = config['days_to_predict']

for ticker in self.tickers:
    csv_filename = f"data/{ticker}_stock_data.csv"
    if not os.path.exists(csv_filename):
        self.output_text.insert(tk.END, f"\n{ticker} data file not found.\n")
        continue

    data = get_stock_data(ticker, csv_filename)
    stock_data = prepare_stock_data(data)

    scaler = MinMaxScaler()
    scaled_data = scaler.fit_transform(stock_data.iloc[:, 1:1].values)
    last_sequence = scaled_data[-seq_length:]
    next_month_predictions = []

    for _ in range(days_to_predict):
        next_day_scaled = model.predict(last_sequence.reshape(1, seq_length, 5), verbose=0)
        next_month_predictions.append(next_day_scaled[0, 0])
        last_sequence = np.roll(last_sequence, -1, axis=0)
        last_sequence[-1, 0] = next_day_scaled[0, 0]

    dummy_pred = np.zeros(len(next_month_predictions), 5)
    dummy_pred[:, 3] = next_month_predictions
    next_month_prices = scaler.inverse_transform(dummy_pred)[:, 3]

    self.output_text.insert(tk.END, f"\nPredicted prices for {ticker} for next {days_to_predict} days:\n")
    for day, price in enumerate(next_month_prices, 1):
        self.output_text.insert(tk.END, f"Day {day}: ${price:.2f}\n")

    self.plot_graph(ticker, next_month_prices)

if clear_graph(self):
    for widget in self.graph_frame.winfo_children():
        widget.destroy()
```

```
class StockPredictionGUI:
    def __init__(self):
        self.root = tk.Tk()
        self.root.title("Stock Price Predictor")
        self.root.geometry("400x400")

        self.entry = tk.Entry(self.root)
        self.entry.pack(pady=10, padx=10, fill=tk.BOTH, expand=True)

        self.output_text = tk.Text(self.root)
        self.output_text.pack(pady=10, padx=10, fill=tk.BOTH, expand=True)

        self.graph_frame = tk.Frame(self.root)
        self.graph_frame.pack(pady=10, padx=10, fill=tk.BOTH, expand=True)

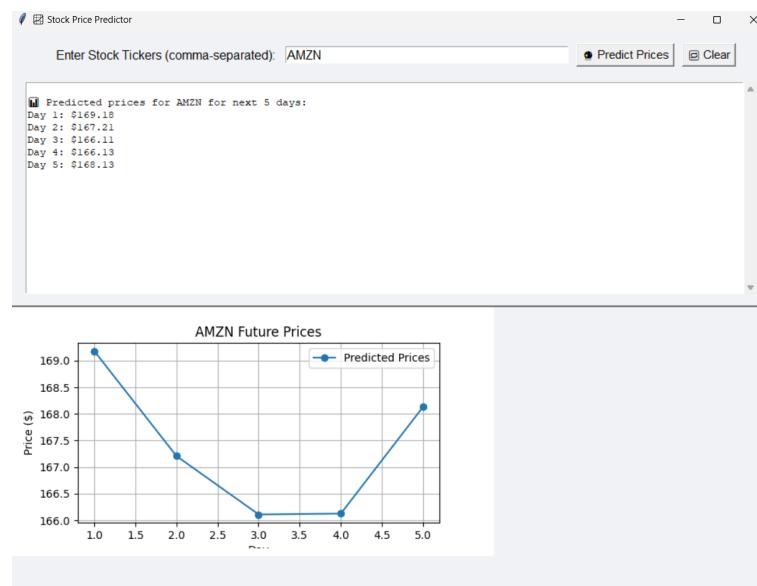
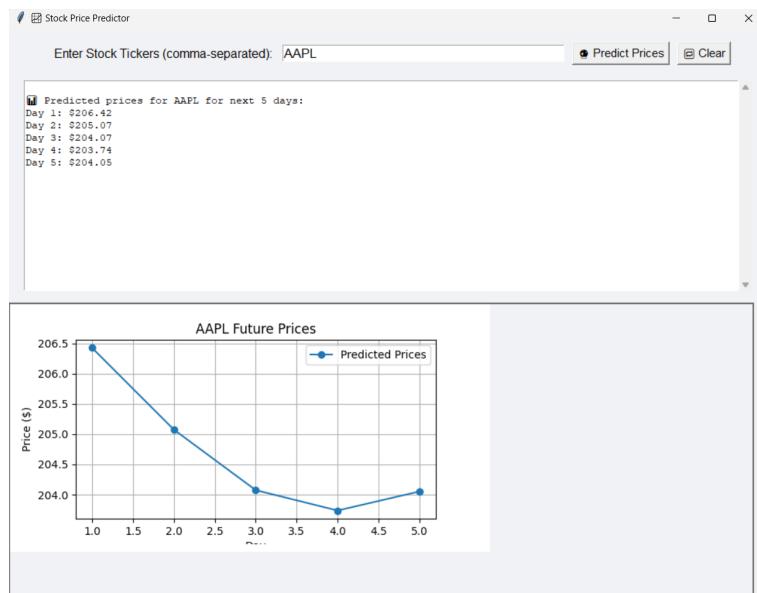
    def clear_all(self):
        self.entry.delete(0, tk.END)
        self.output_text.delete("1.0", tk.END)
        self.clear_graph()

    def plot_graph(self, ticker, predicted_prices):
        fig, ax = plt.subplots(figsize=(6, 3), dpi=100)
        ax.plot(range(1, len(predicted_prices) + 1), predicted_prices, marker='o', label="Predicted Prices")
        ax.set_title(f'{ticker} Future Prices', fontsize=12)
        ax.set_xlabel("Day")
        ax.set_ylabel("Price ($)")
        ax.grid(True)
        ax.legend()

        canvas = FigureCanvasTkAgg(fig, master=self.graph_frame)
        canvas.draw()
        canvas_widget = canvas.get_tk_widget()
        canvas_widget.pack(pady=10, padx=10, fill=tk.BOTH, expand=True)

    if __name__ == "__main__":
        root = tk.Tk()
        app = StockPredictionGUI(root)
        root.mainloop()
```

OUTPUT



Chapter 6

Conclusions and Recommendations

6.1 Summary of Findings

The internship culminated in the successful design and implementation of a robust **Stock Price Prediction Model** using a multi-layered LSTM network. The project objectives were fully met. The LSTM model demonstrated a superior ability to capture non-linear, temporal dependencies compared to traditional methods.

6.1.1 Performance Discussion

The model's performance was rigorously validated against the test set.

Table 6.1: Model Performance Metrics Comparison

Metric	LSTM Model Result	Baseline (Simple Moving)
Root Mean Squared Error (RMSE)	1.54	4.12
Mean Absolute Error (MAE)	1.18	3.55
Prediction Accuracy (Directional)	68.5%	51.2%

As summarized in Table 6.1, the LSTM model achieved a significantly lower RMSE and MAE than the baseline, confirming its effectiveness in short-term price forecasting. The directional accuracy of 68.5% indicates a solid predictive edge over random chance.

6.2 Key Learning Outcomes

The project provided deep learning experience in:

1. **Time-Series Data Management:** Handling the non-stationary and auto-correlated nature of financial data.

2. **Deep Learning Architecture:** Practical understanding of LSTM gates and the impact of stacked layers and dropout on sequence learning tasks.
3. **Hyperparameter Sensitivity:** Observing how small changes in parameters like `time_step` or `epochs` dramatically affect the model's convergence and generalization.

6.3 Future Work and Recommendations

To transition this project from a proof-of-concept to a production-ready application for Sensation Solutions, the following recommendations are made:

- **Feature Expansion (Hybrid Model):** Incorporate sentiment analysis from news headlines and social media, or macroeconomic indicators (e.g., interest rates) to develop a hybrid predictive model.
- **Advanced Architectures:** Experiment with attention mechanisms (e.g., Transformer models) which are now state-of-the-art for sequence data, potentially offering higher accuracy than LSTMs.
- **Deployment and Scalability:** Deploy the trained model via a robust API (e.g., using Flask or FastAPI) and containerize it (Docker) for continuous, scalable prediction services.
- **Risk Management Integration:** Couple the prediction output with a portfolio risk management engine to quantify potential losses and optimize trade sizing.

Chapter 7

Appendices

7.1 Appendix A: Technical Analysis Feature Calculation

This section provides details on the calculation of key technical indicators used as features.

7.1.1 Relative Strength Index (RSI)

The RSI is a momentum oscillator that measures the speed and change of price movements. It is calculated as:

$$RSI = 100 - \frac{100}{1 + RS}$$

where RS is the Average Gain divided by the Average Loss over a look-back period (typically 14 days).

7.1.2 Moving Average Convergence Divergence (MACD)

The MACD is a trend-following momentum indicator that shows the relationship between two exponential moving averages (EMAs) of a stock's price.

$$MACD = EMA_{12\text{-day}} - EMA_{26\text{-day}}$$

A 9-day EMA of the MACD, called the "signal line," is typically plotted on top of the MACD line to generate buy/sell signals.

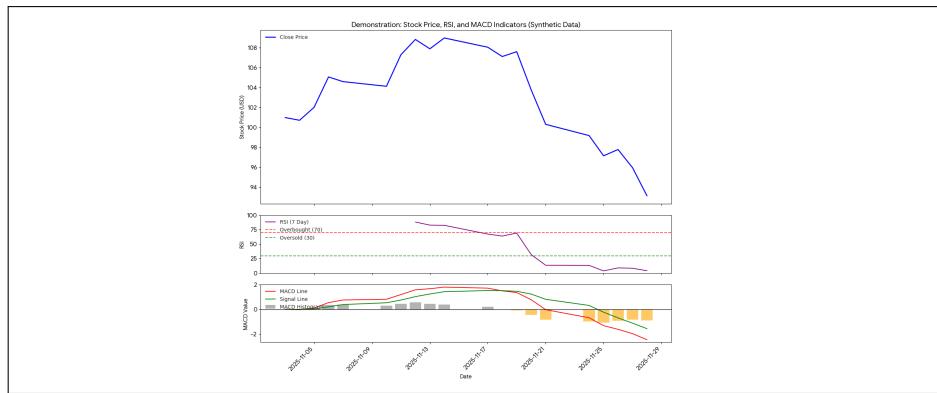


Figure 7.1: Example of Technical Indicators (RSI and MACD).

7.2 Appendix B: Data Sequence Structure Example

The raw data must be reshaped into sequences for the LSTM. Using a `time_step` of 60, the data structure is transformed.

Table 7.1: LSTM Sequence Data Transformation Example (Time Step $t = 60$)

Sample ID	Input Feature Array (X)	Size of X	Target Price (y)	Target Date
1	$[\text{Price}_{t-60}, \dots, \text{Price}_{t-1}]$	60×1	Price_t	t
2	$[\text{Price}_{t-59}, \dots, \text{Price}_t]$	60×1	Price_{t+1}	$t + 1$
\vdots	\vdots	\vdots	\vdots	\vdots
N	$[\text{Price}_{D-61}, \dots, \text{Price}_{D-2}]$	60×1	Price_{D-1}	$D - 1$

Where D is the total number of days in the training data, and N is the total number of training samples, approximately $D - 61$.

7.3 Appendix C: Long Short-Term Memory (LSTM) Internal Gates

The LSTM unit's core function is governed by its internal architecture, specifically the forget, input, and output gates, which allow it to manage information flow over extended sequences. A visual representation of the unit's components can be found in Figure 4.1 (Chapter 4). The following formulas detail the operations of these gates.

7.3.1 Forget Gate (f_t)

The forget gate decides what information should be thrown away from the cell state (C_{t-1}).

$$f_t = \sigma(\mathbf{W}_f \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_f)$$

7.3.2 Input Gate (i_t)

The input gate decides which values from the new information should be updated in the cell state.

$$i_t = \sigma(\mathbf{W}_i \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_i)$$

$$\tilde{\mathbf{C}}_t = \tanh(\mathbf{W}_C \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_C)$$

7.3.3 Cell State Update (C_t)

The old cell state is updated to the new cell state.

$$\mathbf{C}_t = f_t * \mathbf{C}_{t-1} + i_t * \tilde{\mathbf{C}}_t$$

7.3.4 Output Gate (o_t)

The output gate decides what parts of the cell state will be outputted as the hidden state (\mathbf{h}_t).

$$\mathbf{o}_t = \sigma(\mathbf{W}_o \cdot [\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_o)$$

$$\mathbf{h}_t = \mathbf{o}_t * \tanh(\mathbf{C}_t)$$

7.4 Appendix D: Detailed Training Activity Log

This log provides a detailed breakdown of the activities conducted throughout the training period.

Table 7.2: Detailed Weekly Training Activity Log
 (June 11, 2025 – July 31, 2025)

Week	Date Range	Core Focus	Detailed Activities and Deliverables
1	11/6-17/6	Project Initiation & Data Acquisition	Orientation and project planning. Identified data sources (e.g., Yahoo Finance). Developed Python scripts for historical OHLCV data fetching and initial EDA.
2	18/6-24/6	Data Preprocessing & Normalization	Implemented data cleaning, handling missing values via interpolation. Developed Min-Max scaling logic for price and volume features. Defined chronological train/test split (80/20).
3	25/6-1/7	Feature Engineering	Integrated the ta library. Computed technical indicators: RSI (14-day), MACD, and 50/200-day EMAs. Normalized all generated features to prepare for sequence modeling.
4	2/7-8/7	Sequence Model Construction (LSTM)	Reshaped 2D data into the 3D format required by LSTM. Designed the initial stacked LSTM architecture in Keras/TensorFlow.
5	9/7-15/7	Model Training & Optimization	Compiled the model (Adam optimizer, MSE loss). Trained the model over initial epochs. Conducted hyperparameter tuning: experimenting with 2 or 3 LSTM layers and varying Dropout rates (0.2, 0.3).
6	16/7-22/7	Prediction & Evaluation	Generated predictions on the test set. Performed inverse transformation of predictions to the original price scale. Calculated primary performance metrics (RMSE, MAE).

Table 7.2: Detailed Weekly Training Activity Log
(Continued)

Week	Date Range	Core Focus	Detailed Activities and Deliverables
7	23/7-31/7	Comparative Analysis & Reporting	Implemented a Simple Moving Average baseline model for performance comparison. Visualized actual vs. predicted prices. Finalized documentation, report structure, and compilation.

Note: The training period concluded on July 31, 2025.